

# README

HPA Demo Project (Dynamic Scale Up / Scale Down) Goal Demonstrate

HorizontalPodAutoscaler behavior in a practical way where: - replicas scale **up** automatically under CPU pressure - replicas scale **down** automatically after load is removed This project is intentionally simple and runnable on local labs (Minikube/KIND) or any Kubernetes cluster with metrics support. Project structure text hpa-demo/ README.md manifests/ 00-namespace.yaml 01-deployment-service.yaml 02-hpa.yaml 03-load-generator.yaml How this demo works App deployment: hpa-web (CPU request defined: 100m) HPA target: average CPU utilization = 50% HPA limits: min 1, max 10 Load generator: continuously sends HTTP traffic to service hpa-web-svc Because HPA uses **utilization relative to CPU request**, setting resources.requests.cpu is required.

Prerequisites Running Kubernetes cluster kubectl configured Metrics pipeline available (typically metrics-server) Check metrics quickly: bash kubectl top nodes kubectl top pods -A If these commands fail, install metrics-server first (method depends on your cluster distro). If metrics-server is not running (HPA shows unknown metrics) Check current status: bash kubectl get apiservice v1beta1.metrics.k8s.io kubectl -n kube-system get deploy,pods | findstr metrics-server Install metrics-server (upstream manifest): bash kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml For local/lab clusters (KIND, some Minikube setups), kubelet cert validation can block metrics collection. Add these args to metrics-server deployment if logs show TLS / x509 / scrape errors: bash kubectl -n kube-system edit deployment metrics-server Under spec.template.spec.containers[0].args, ensure these flags exist: text --kubelet-insecure-tls

--kubelet-preferred-address-types=InternalIP,Hostname,ExternalIP Wait for rollout and verify: bash kubectl -n kube-system rollout status deploy/metrics-server kubectl get apiservice v1beta1.metrics.k8s.io kubectl top nodes kubectl top pods -A If kubectl top works, HPA should start receiving metrics within ~30-90 seconds. Step 1: Deploy app + HPA bash kubectl apply -f hpa-demo/manifests/00-namespace.yaml kubectl apply -f hpa-demo/manifests/01-deployment-service.yaml kubectl apply -f hpa-demo/manifests/02-hpa.yaml Verify baseline: bash kubectl get deploy,hpa,pods -n hpa-demo kubectl describe hpa hpa-web -n hpa-demo

Expected initial state: - deployment replicas around 1 - HPA target shown as 50% Step 2: Start load and observe scale up Apply load generator: bash kubectl apply -f

hpa-demo/manifests/03-load-generator.yaml Watch autoscaling behavior: bash kubectl get hpa hpa-web -n hpa-demo -w In another terminal: bash kubectl get deploy hpa-web -n hpa-demo -w You should see: - TARGETS in HPA rise above 50% - REPLICAS increase dynamically (e.g., 1 -> 2 -> 4 ... depending on cluster capacity) Step 3: Stop load and observe scale down Remove load generator: bash kubectl delete -f

hpa-demo/manifests/03-load-generator.yaml Keep watching HPA/deployment for 1-3 minutes. You should see: - CPU utilization fall - replicas reduce gradually back toward

minReplicas: 1 This project sets scaleDown.stabilizationWindowSeconds: 30 for

quicker classroom/demo feedback. Useful debug commands bash kubectl top pods -n hpa-demo kubectl describe hpa hpa-web -n hpa-demo kubectl get events -n hpa-demo --sort-by=.lastTimestamp kubectl logs -n hpa-demo

deploy/load-generator --tail=20 Practical notes (important) **HPA is reactive, not instant**

It takes some time to collect metrics and apply scaling decisions. **CPU request is mandatory for CPU utilization target** Without resources.requests.cpu, utilization-based CPU scaling does not work as intended. **Scale down is intentionally slower in real systems** Avoids oscillation/flapping. This demo uses a shorter window (30s) for visible lab behavior. **Cluster limits matter** If nodes are resource-constrained, HPA may request replicas that remain Pending. **HPA does not replace Cluster Autoscaler** HPA scales pods; Cluster Autoscaler scales nodes.

Cleanup bash kubectl delete -f hpa-demo/manifests/03-load-generator.yaml --ignore-not-found=true kubectl delete -f hpa-demo/manifests/02-hpa.yaml kubectl delete -f hpa-demo/manifests/01-deployment-service.yaml kubectl

```
delete -f hpa-demo/manifests/00-namespace.yaml
```